


Cmod A7-35T Memory Specification

Platform: Cmod A7-35T RISC-V MCU

Scope: memory hierarchy, address map, caches, measured access latencies, placement guidance

The MCU provides three kinds of memory: single-cycle tightly-coupled RAM, cached external SRAM, and QSPI flash storage. Ordinary code and data reside in the large SRAM; latency-critical code and data can be placed in memories with fixed single-cycle timing.


1. Memory Hierarchy Overview



Level	Memory	Size	Access	Role
1	ITCM + DTCM (tightly-coupled RAM)	32 KB + 64 KB	1 cycle, fixed	interrupt handlers, stack, hot data
2	I-Cache + D-Cache	16 KB + 16 KB	adds ≈0 cycles on hit	transparent acceleration of SRAM
3	External SRAM	512 KB	≈26 cycles/word on miss	application code, data, heap
—	QSPI flash	4 MB	not directly addressable	firmware storage, read at boot

The tightly-coupled memories sit on a dedicated dual-port memory bus with guaranteed single-cycle access; they are never cached because they are already as fast as a cache hit. The external SRAM is the application's main memory; it is reached through the caches, so hot code and data run at tightly-coupled speed while the total space is 512 KB. The QSPI flash has no CPU address window: firmware is stored in flash and executed from SRAM, and the bootloader performs the copy at every power-on.

2. Memory Map



Address range	Size	Region	Access
0x0000_0000 – 0x0000_7FFF	32 KB	Bootloader (reserved)	1 cycle
0x0000_8000 – 0x0000_FFFF	32 KB	ITCM — application fast code	1 cycle
0x0001_0000 – 0x0001_FFFF	64 KB	DTCM — fast data; stack grows down from 0x0002_0000	1 cycle
0x4000_0000 – 0x40FF_FFFF	—	Peripheral registers	uncached
0x6000_0000 – 0x6007_FFFF	512 KB	External SRAM — application code / data / heap	cached

Address decoding is exact: an access outside these ranges raises a bus error instead of aliasing onto real memory, so a stray pointer is detected rather than silently corrupting data. The QSPI flash does not appear in

this map; it is accessed only through the flash controller's registers. Peripheral registers and base addresses are documented in the IP peripheral reference.

3. Caches

Item	Value
Instruction cache	16 KB, 32-byte lines
Data cache	16 KB, 32-byte lines
Write policy	write-through — every store is forwarded to SRAM
Cached range	exactly the SRAM: 0x6000_0000 – 0x6007_FFFF
Software control	none — always on for the cached range

Only the SRAM range is cached. The tightly-coupled memories are already single-cycle, and peripheral registers must observe every access, so neither is ever cached. There is no cache enable/disable or flush to manage in normal application code.

The write-through policy has two practical consequences:

- Reads benefit fully from the cache: a hit costs the same as a tightly-coupled RAM access.
- Writes always pay the SRAM latency, even on a cache hit. Frequently written data (counters, buffers filled in interrupt handlers) therefore belongs in the DTCM.

4. Measured Access Latencies

Measured on the board at 100 MHz with an optimization-enabled word-copy loop. Values are cycles per 32-bit word and include the loop's own ≈ 7 -cycle overhead; the difference between rows is the memory's contribution.

Access	Cycles/word	Notes
Tightly-coupled RAM read	≈ 7.2	equals the bare loop; the memory adds no cycles
SRAM read, cache hit	≈ 7.2	a hit is as fast as tightly-coupled RAM
SRAM read, cache miss	≈ 26	≈ 151 cycles to fill one 32-byte line
SRAM write (write-through)	≈ 23.6	every store pays the SRAM latency

Note: Measure with compiler optimization enabled. An unoptimized build inflates the loop overhead to the point where cached and uncached accesses look almost identical; the memory effects are hidden by the extra instructions.

5. Memory Placement Guidance

What	Put it in	How (course template)
Interrupt handlers, timing-critical loops	ITCM	tag the function <code>ITCM_FUNC</code>

What	Put it in	How (course template)
Frequently-written data, DMA-like buffers	DTCM	tag the variable <code>DTCM_DATA</code>
Stack	DTCM	placed there by the template's linker script
Everything else — code, constants, data, heap	SRAM	default
Firmware image	QSPI flash	written by the deployment step

As a general rule, place code and data in SRAM by default and move them to the TCMs when timing must be predictable rather than merely fast. A function in SRAM usually runs at cache speed, but its worst case (a cold cache line) costs three to four times more; for an interrupt handler this spread appears directly as latency jitter. In the tightly-coupled memories the worst case equals the best case.

At power-on the bootloader copies the application image from flash into SRAM (see the Standalone Boot Mode guide); the template's `mcu_init()` then copies `ITCM_FUNC` code out of the SRAM image into the ITCM and zeroes the DTCM. For this reason the template requires `mcu_init();` to remain the first line of `main()`: nothing placed in the tightly-coupled memories works before it runs. The memory tour printed at boot lists each region's live addresses and serves as a check that the layout is in effect.